



Embedded Operating System

Mutual exclusion and deadlocks

Graziano Pravadelli

Mutual exclusion

Mutal exclusion

Assumptions:

- n processes; each process P_i running on a different processor
- Each process has a critical section to be executed in mutual exclusion

Requirements

- If P_i it is executing its own critical section, there is no other process P_j which is executing its own critical section

Two types of approaches

- Centralized
- Distributed


Centralized approach

A process takes care of coordinating access to critical section (CS)


- A process that wants to enter the CS sends a **request** message to the coordinator
- The coordinator decides which process enters the CS and sends a **reply** message
- The process that receives the **reply** enters the CS
- When it exits from the CS, it sends a **release** message to the coordinator

Three messages for each access in the CS


Centralized approach – Example



(a)



(b)



(c)



Centralized approach – Pros & Cons

Advantages

- Simple (3 messages for each request)
- It guarantees the three conditions (mutual exclusion, progress, limited waiting) without deadlock/starvation
- Suitable for all kinds of resources

Disadvantages

- The coordinator is a critical point
 - Fault tolerance (if it crashes?)
 - Performances (possible bottleneck)

Distributed approach [Ricart & Agrawala 81]

- Based on broadcast and on the total ordering of events in the system
- Given a couple of events the algorithm uniquely defines which has occurred before
- We can use Lamport's timestamps

Distributed approach

When a process P_i wants to enter the CS

- It generates a timestamp TS
- It sends *request* (P_i, CS_k, TS) to all processes
- The send of messages is reliable (ACK for each message)

When a process P_j receives a *request* message

- It can
 - send a *reply* (OK) message
 - or put the request in a queue

When a process P_i receives a *reply* (OK) message from all other processes

- It can enter the CS

When P_i leaves the CS

- send a *reply* (OK) message to all queued requests




Distributed approach

Does P_j send a *reply* message or not?

When P_j receives *request*(P_i, CS_k, TS)

- If P_j is in CS, it enqueues the *reply* (OK) to P_i
- If P_j is not in CS and it does not want to enter CS, it sends *reply* (OK) to P_i
- If P_j wants to enter CS, it compares the timestamp of its own request TS_j with TS :
 - If $TS_j > TS$, it sends the *reply* (OK) to P_i (P_i has requested before P_j)
 - Otherwise, it enqueues the *reply* (OK)

Distributed approach – Example



Distributed approach – Pros & Cons

Advantages

- Mutual exclusion guaranteed without deadlock/starvation
- $2(n-1)$ for each access in the CS
- No a single critical point

Disadvantages

- n critical points
 - Crash of a process = No access to the CS
- n bottlenecks
- Need group communication primitives

Variant

- Instead of enqueueing messages, explicit rejection message
- It acts as ACK and solves the problem of n critical points

Token-ring algorithm



Token originally is in process 0



It pass through point-to-point messages from process i to $i+1 \text{ (module } N)$




When a process receives the token

if it wants to enter the CS, it enters the CS. When it leaves the CS, it passes the token to the next process

if it does not want to enter the CS, it immediately passes the token

Token-ring algorithm – Example

- It assumes a specific logical organization (ring) of processes
- Based on the rotation of a token between various processes



Token-ring algorithm – Pros & Cons


Advantages

- No deadlock/starvation
- The impact of node failure is limited
 - In the event of a crash, it is sufficient to "bypass" (updating the network structure)
 - We notice the crash if the ACK of the token does not arrive

Disadvantages

- Loss of the token
 - How to detect it (long CS vs. lost token)?

Deadlock



Definition



Processes use resources



Sequence of use

Request
[Wait]
Use
Release



DEADLOCK

A set of processes is in *deadlock* when each process is waiting for an event that can be caused by a process of the same set

Required conditions

Mutual exclusion

- At least one resource must be non-shareable

Hold and Wait

- There must be a process that holds a resource and expects to acquire another resource, held by another

No preemption

- Resources can only be released “voluntarily” by the process that is using them

Circular waiting

- There must be a set of processes that cyclically wait the release of a resource

Deadlock management

- Same options of centralized systems
 - Ostrich algorithm
 - **Static prevention**
 - **Avoid circular waiting**
 - **Distributed**
 - Dynamic prevention based on resource allocation
 - **Detection and recovery**
 - **Centralized**
 - **Distributed**

RAG model

Resource allocation graph (RAG) $G(V,E)$

- V = nodes
- E = edges

Nodes


- Circles = processes (CPU, I/O, memory)
- Rectangles = resources
 - In the rectangles there are as many "•" as the instances of the corresponding resource

Edges

- From processes to resources: process requires resource
- From resources to processes: process holds resource

Detection – Centralized approach


- Emulation of the not distributed case
- A coordinator maintains a global RAG
 - Union of individual RAGs of individual hosts
 - If it finds a cycle, it kills a process



Detection – Centralized approach

- How and when to build the global RAG? (when send it to the administrator?)
 - After each addition/removal of an arc in the local graphs
 - Periodically
 - At the request of the coordinator

False deadlocks – Example



- B releases R and requests T (legal operation)
- What is the order of the messages?
 - Msg1: Host 2 announces to the coordinator that B releases R
 - Msg2: Host 1 announces to the coordinator that B requests T
- If messages arrive in opposite order → false deadlock

False deadlocks – Solution

- Global graph must be updated independently from the order in which messages are received
 - Association of (Lamport) timestamps to modifications of local graphs
 - When the coordinator receives a change with timestamp T , the coordinator itself urges the sending of any other changes in progress with timestamp smaller than T

Detection – Distributed approach

[Chandy&Misra&Haas 83]

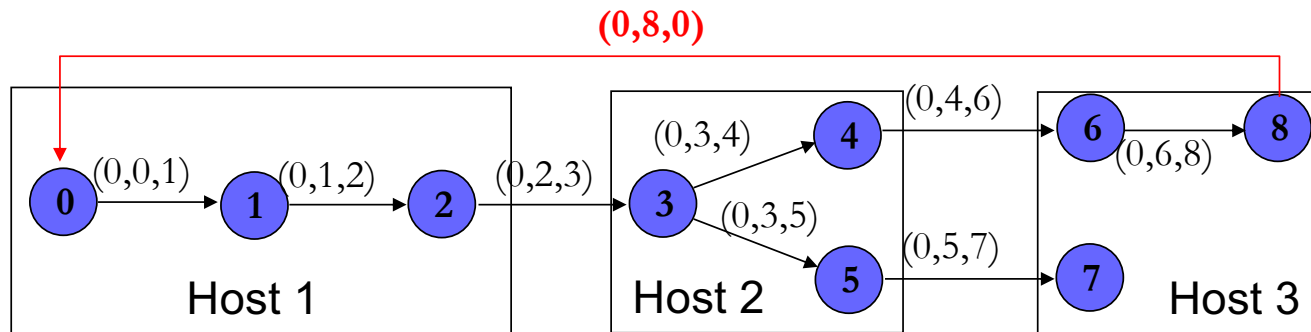
Invoked when a process waits for some resource

- It sends probe messages to processes that hold the resource
 - **Message = (id blocked proc, id sender proc, id destination proc)**
- The probe is propagated to processes that are also waiting
 - If the receiver is waiting, it forwards the probe by updating the sender and receiver
- If the message reaches the sender → deadlock

Recovery


- Killing the probe initiator → possible many kills if many processes start the algorithm at the same time
- Killing the process with the higher id

Distributed approach – Example



Static prevention based on resource ordering

- Preventing the circular waiting condition by definition of a global order among resources in the system
 - E.g.: assigning a single number to each resource
 - A process can request a resource with a number n only if it does not hold resources with numbers greater than n
- Simple and low-cost implementation



Static prevention based on timestamp

Timestamp used as a priority value for each process

Use of timestamps to regulate the wait of a process

- Non-preemptive scheme (**wait-die**)
- Preemptive scheme (**wound-wait**)

The use of timestamps prevents from starvation

Wait-Die scheme

- If a process P_i needs a resource which is held by a process P_j
 - if $TS(P_i) < TS(P_j) \Rightarrow P_i$ goes on wait (wait)
 - if $TS(P_i) > TS(P_j) \Rightarrow$ rollback of P_i (die)

Wait-die – Example

wait



TS = 10



TS = 99

die



TS = 110



TS = 10



**Always in the direction
of increasing timestamps**

Wound-Wait scheme

- If a process P_i needs a resource which is held by another process P_j
 - if $TS(P_i) < TS(P_j) \Rightarrow P_j$ is wounded and killed (wound)
 - if $TS(P_i) > TS(P_j) \Rightarrow$ it goes on wait (wait)

Wound-wait – Example

wound



TS = 10



TS = 110

wait



TS = 99



TS = 10



**Always in the direction
of decreasing timestamps**

Static prevention based on timestamp

- Wait-die tends to kill many young processes
 - The old respectful man waits
 - The presumptuous young man is repeatedly killed
- Wound-wait kills only once
 - The wise old man pre-empts the young man
 - The young man re-starts and waits patiently